

[Get started](#)[Open in app](#)[Follow](#)

578K Followers



PyTorch vs TensorFlow — spotting the difference



Kirill Dubovikov · Jun 20, 2017 · 9 min read



In this post I want to explore some of the key similarities and differences between two popular deep learning frameworks: PyTorch and TensorFlow. Why those two and not the others? There are many deep learning frameworks and many of them are viable tools, I chose those two just because I was interested in comparing them specifically.

[Get started](#)[Open in app](#)

and production needs. Its closed-source predecessor is called DistBelief.

PyTorch is a cousin of lua-based Torch framework which was developed and used at Facebook. However, PyTorch is not a simple set of wrappers to support popular language, it was rewritten and tailored to be fast and feel native.

The best way to compare two frameworks is to code something up in both of them. I've written a companion jupyter notebook for this post and you can [get it here](#). All code will be provided in the post too.

First, let's code a simple approximator for the following function in both frameworks:

$$f(x) = x^\phi$$

We will try to find unknown parameter ϕ given data x and function values $f(x)$. Yes, using stochastic gradient descent for this is an overkill and analytical solution may be found easily, but this problem will serve our purpose well as a simple example.

We will solve this with PyTorch first:

```
1  import torch
2  from torch.autograd import Variable
3  import numpy as np
4
5  def rmse(y, y_hat):
6      """Compute root mean squared error"""
7      return torch.sqrt(torch.mean((y - y_hat).pow(2).sum()))
8
9  def forward(x, e):
10     """Forward pass for our fuction"""
11     return x.pow(e.repeat(x.size(0)))
12
13  # Let's define some settings
14  n = 100 # number of examples
```

Get started

Open in app



```
17
18 # Model definition
19 x = Variable(torch.rand(n) * 10, requires_grad=False)
20
21 # Model parameter and it's true value
22 exp = Variable(torch.FloatTensor([target_exp]), requires_grad=False)
23 exp_hat = Variable(torch.FloatTensor([4]), requires_grad=True) # just some starting value, could
24 y = forward(x, exp)
25
26 # a couple of buffers to hold parameter and loss history
27 loss_history = []
28 exp_history = []
29
30 # Training loop
31 for i in range(0, 200):
32     print("Iteration %d" % i)
33
34     # Compute current estimate
35     y_hat = forward(x, exp_hat)
36
37     # Calculate loss function
38     loss = rmse(y, y_hat)
39
40     # Do some recordings for plots
41     loss_history.append(loss.data[0])
42     exp_history.append(y_hat.data[0])
43
44     # Compute gradients
45     loss.backward()
46
47     print("loss = %s" % loss.data[0])
48     print("exp = %s" % exp_hat.data[0])
49
50     # Update model parameters
51     exp_hat.data -= learning_rate * exp_hat.grad.data
52     exp_hat.grad.data.zero_()
```

pytorch.py hosted with ❤ by GitHub

[view raw](#)

If you have some experience in deep learning frameworks you may have noticed that we are implementing gradient descent by hand. Not very convenient, huh? Gladly, PyTorch

Get started

Open in app

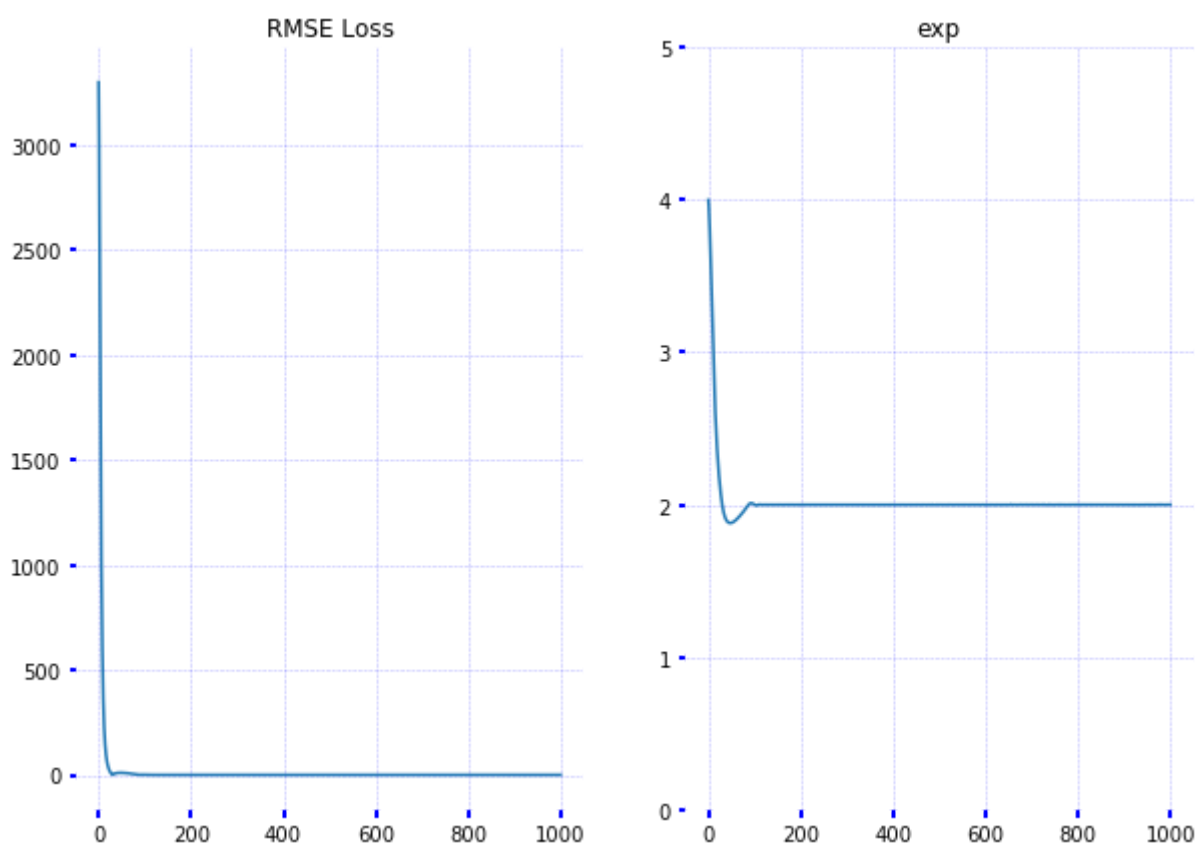


```
1  import torch
2  from torch.autograd import Variable
3  import numpy as np
4
5  def rmse(y, y_hat):
6      """Compute root mean squared error"""
7      return torch.sqrt(torch.mean((y - y_hat).pow(2)))
8
9  def forward(x, e):
10     """Forward pass for our fuction"""
11     return x.pow(e.repeat(x.size(0)))
12
13  # Let's define some settings
14  n = 1000 # number of examples
15  learning_rate = 5e-10
16
17  # Model definition
18  x = Variable(torch.rand(n) * 10, requires_grad=False)
19  y = forward(x, exp)
20
21  # Model parameters
22  exp = Variable(torch.FloatTensor([2.0]), requires_grad=False)
23  exp_hat = Variable(torch.FloatTensor([4]), requires_grad=True)
24
25  # Optimizer (NEW)
26  opt = torch.optim.SGD([exp_hat], lr=learning_rate, momentum=0.9)
27
28  loss_history = []
29  exp_history = []
30
31  # Training loop
32  for i in range(0, 10000):
33      opt.zero_grad()
34      print("Iteration %d" % i)
35
36      # Compute current estimate
37      y_hat = forward(x, exp_hat)
38
39      # Calculate loss function
40      loss = rmse(y, y_hat)
41
```

[Get started](#)[Open in app](#)

```
44     exp_history.append(y_hat.data[0])
45
46     # Update model parameters
47     loss.backward()
48     opt.step()
49
50     print("loss = %s" % loss.data[0])
51     print("exp = %s" % exp_hat.data[0])
```

pytorch_optim.py hosted with ❤ by GitHub

[view raw](#)

Loss function and exponent plots for PyTorch

As you can see, we quickly inferred true exponent from training data. And now let's go on with TensorFlow:

```
1  import tensorflow as tf
2
3  def rmse(y, y_hat):
4      """Compute root mean squared error"""
```

Get started

Open in app

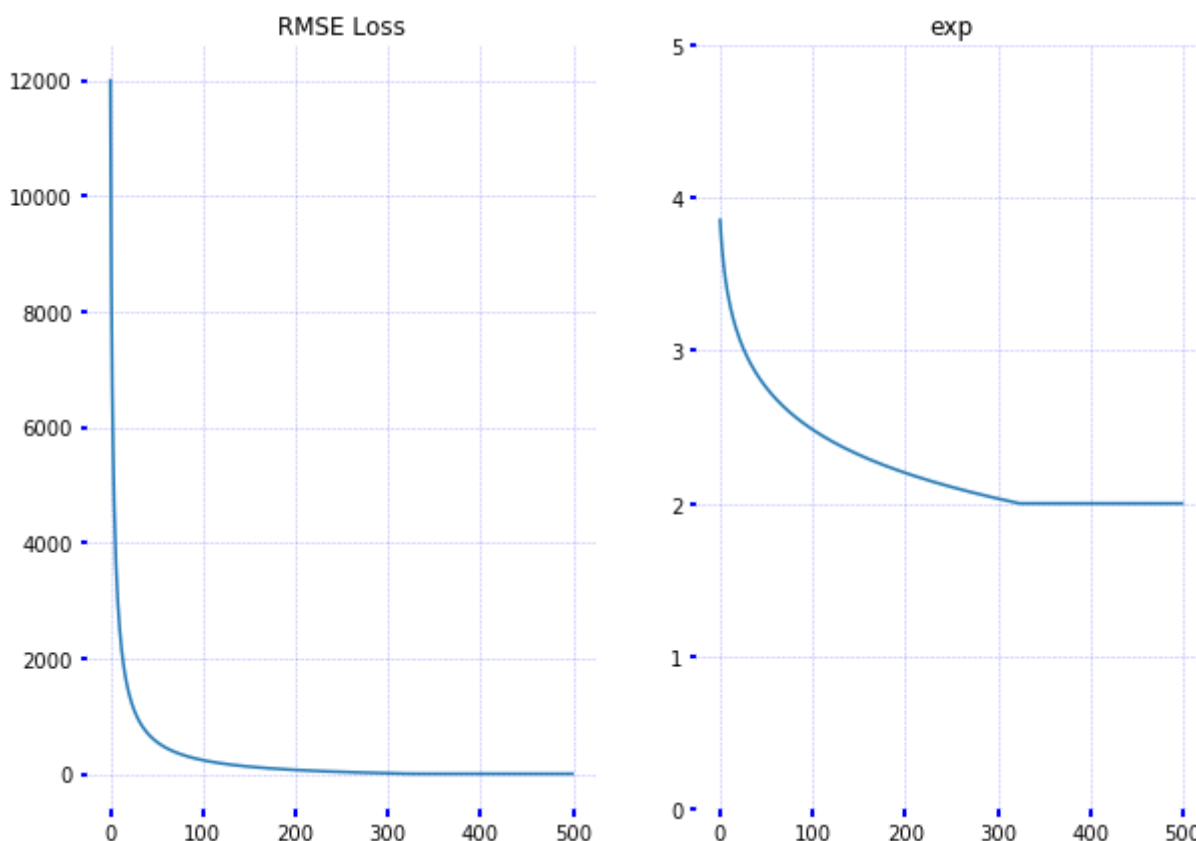


```
8     """Forward pass for our fuction"""
9     # tensorflow has automatic broadcasting
10    # so we do not need to reshape e manually
11    return tf.pow(x, e)
12
13    n = 100 # number of examples
14    learning_rate = 5e-6
15
16    # Placeholders for data
17    x = tf.placeholder(tf.float32)
18    y = tf.placeholder(tf.float32)
19
20    # Model parameters
21    exp = tf.constant(2.0)
22    exp_hat = tf.Variable(4.0, name='exp_hat')
23
24    # Model definition
25    y_hat = forward(x, exp_hat)
26
27    # Optimizer
28    loss = rmse(y, y_hat)
29    opt = tf.train.GradientDescentOptimizer(learning_rate)
30
31    # We will run this operation to perform a single training step,
32    # e.g. opt.step() in Pytorch.
33    # Execution of this operation will also update model parameters
34    train_op = opt.minimize(loss)
35
36    # Let's generate some training data
37    x_train = np.random.rand(n) + 10
38    y_train = x_train ** 2
39
40    loss_history = []
41    exp_history = []
42
43    # First, we need to create a Tensorflow session object
44    with tf.Session() as sess:
45
46        # Initialize all defined variables
47        tf.global_variables_initializer().run()
48
49        # Training loop
```

[Get started](#)[Open in app](#)

```
53     curr_loss, curr_exp, _ = sess.run([loss, exp_hat, train_op], feed_dict={x: x_train, y: y_train})
54
55     print("loss = %s" % curr_loss)
56     print("exp = %s" % curr_exp)
57
58     # Do some recordings for plots
59     loss_history.append(curr_loss)
60     exp_history.append(curr_exp)
```

tf_example_1.py hosted with ❤ by GitHub

[view raw](#)

Loss function and exponent plots for TensorFlow

As you can see, implementation in TensorFlow works too (surprisingly 🤔). It took more iterations to recover the exponent, but I am sure that the cause is I did not fiddle with optimiser's parameters enough to reach comparable results.

Now we are ready to explore some differences.

[Get started](#)[Open in app](#)

professionals. The framework is well documented and if the documentation will not suffice there are many extremely well-written tutorials on the internet. You can find hundreds of implemented and trained models on github, start here.

PyTorch is relatively new compared to its competitor (and is still in beta), but it is quickly getting its momentum. Documentation and official tutorials are also nice. PyTorch also include several implementations of popular computer vision architectures which are super-easy to use.

Difference #1 — dynamic vs static graph definition

Both frameworks operate on tensors and view any model as a directed acyclic graph (DAG), but they differ drastically on how you can define them.

TensorFlow follows ‘data as code and code is data’ idiom. In TensorFlow you define graph statically before a model can run. All communication with outer world is performed via `tf.Session` object and `tf.Placeholder` which are tensors that will be substituted by external data at runtime.

In PyTorch things are way more imperative and dynamic: you can define, change and execute nodes as you go, no special session interfaces or placeholders. Overall, the framework is more tightly integrated with Python language and feels more native most of the times. When you write in TensorFlow sometimes you feel that your model is behind a brick wall with several tiny holes to communicate over. Anyways, this still sounds like a matter of taste more or less.

However, those approaches differ not only in a software engineering perspective: there are several dynamic neural network architectures that can benefit from the dynamic approach. Recall RNNs: with static graphs, the input sequence length will stay constant. This means that if you develop a sentiment analysis model for English sentences you must fix the sentence length to some maximum value and pad all smaller sequences with zeros. Not too convenient, huh. And you will get more problems in the domain of recursive RNNs and tree-RNNs. Currently Tensorflow has limited support for dynamic inputs via Tensorflow Fold. PyTorch has it by-default.

Get started

Open in app



Python debugging tools such as pdb, ipdb, PyCharm debugger or old trusty print statements.

This is not the case with TensorFlow. You have an option to use a special tool called `tfdbg` which allows to evaluate tensorflow expressions at runtime and browse all tensors and operations in session scope. Of course, you won't be able to debug any python code with it, so it will be necessary to use pdb separately.

```

--- run-start: run #1: 1 fetch (accuracy/accuracy/Mean:0); 2 feeds -----
| <-- --> | run_info
| run | invoke stepper | exit |
UP

TTTTT FFFF DDD BBBB GGG
TT  F  D D B  B G
TT  FFF D D BBBB G GG
TT  F  D D B  B G G
TT  F   DD BBBB GGG

=====
Session.run() call #1:

Fetch(es):
  accuracy/accuracy/Mean:0

Feed dict(s):
  input/x-input:0
  input/y-input:0

=====

Select one of the following commands to proceed ---->
  run:
    Execute the run() call with debug tensor-watching
  run -n:
    Execute the run() call without debug tensor-watching
DN
--- Scroll (PgDn): 0.00% ----- Mouse: ON --
tfdbg>

```

Difference #3 — Visualization

Tensorboard is awesome when it comes to visualization 😎. This tool comes with TensorFlow and it is very useful for debugging and comparison of different training runs. For example, consider you trained a model, then tuned some hyperparameters and trained it again. Both runs can be displayed at Tensorboard simultaneously to indicate possible differences. Tensorboard can:

- Display model graph

[Get started](#)[Open in app](#)

Visualize distributions and histograms

- Visualize images
- Visualize embeddings
- Play audio

Tensorboard can display various summaries which can be collected via `tf.summary` module. We will define summary operations for our toy exponent example and use `tf.summary.FileWriter` to save them to disk.

```
1  import tensorflow as tf
2  import numpy as np
3
4  def rmse(y, y_hat):
5      """Compute root mean squared error"""
6      return tf.sqrt(tf.reduce_mean(tf.square((y - y_hat))))
7
8  def forward(x, e):
9      """Forward pass for our fuction"""
10     # tensorflow has automatic broadcasting
11     # so we do not need to reshape e manually
12     return tf.pow(x, e)
13
14  n = 100 # number of examples
15  learning_rate = 5e-6
16
17  # Placeholders for data
18  x = tf.placeholder(tf.float32)
19  y = tf.placeholder(tf.float32)
20
21  # Model parameters
22  exp = tf.constant(2.0)
23  exp_hat = tf.Variable(4.0, name='exp_hat')
24
25  # Model definition
26  y_hat = forward(x, exp_hat)
27
28  # Optimizer
29  loss = rmse(y, y_hat)
```

[Get started](#)[Open in app](#)

```
33 loss_summary = tf.summary.scalar("loss", loss)
34 exp_summary = tf.summary.scalar("exp", exp_hat)
35 all_summaries = tf.summary.merge_all()
36
37 # We will run this operation to perform a single training step,
38 # e.g. opt.step() in Pytorch.
39 # Execution of this operation will also update model parameters
40 train_op = opt.minimize(loss)
41
42 # Let's generate some training data
43 x_train = np.random.rand(n) + 10
44 y_train = x_train ** 2
45
46 loss_history = []
47 exp_history = []
48
49 # First, we need to create a Tensorflow session object
50 with tf.Session() as sess:
51
52     # Initialize all defined variables
53     tf.global_variables_initializer().run()
54
55     summary_writer = tf.summary.FileWriter('./tensorboard', sess.graph)
56
57     # Training loop
58     for i in range(0, 500):
59         print("Iteration %d" % i)
60         # Run a single training step
61         summaries, curr_loss, curr_exp, _ = sess.run([all_summaries, loss, exp_hat, train_op], f
62
63         print("loss = %s" % curr_loss)
64         print("exp = %s" % curr_exp)
65
66         # Do some recordings for plots
67         loss_history.append(curr_loss)
68         exp_history.append(curr_exp)
69
70         summary_writer.add_summary(summaries, i)
```

tensorboard.py hosted with ❤ by GitHub

[view raw](#)

[Get started](#)[Open in app](#)

Tensorboard competitor from the PyTorch side is [visdom](#). It is not as feature-complete, but a bit more convenient to use. Also, [integrations](#) with Tensorboard do exist. Also, you are free to use standard plotting tools — [matplotlib](#) and [seaborn](#).

Difference #4 — Deployment

If we start talking about deployment TensorFlow is a clear winner for now: it has [TensorFlow Serving](#) which is a framework to deploy your models on a specialized gRPC server. Mobile is also supported.

When we switch back to PyTorch we may use [Flask](#) or another alternative to code up a REST API on top of the model. This could be done with TensorFlow models as well if gRPC is not a good match for your usecase. However, TensorFlow Serving may be a better option if performance is a concern.

Tensorflow also supports [distributed training](#) which PyTorch lacks for now.

Difference #5 — Data Parallelism

One of the biggest features that distinguish PyTorch from TensorFlow is [declarative data parallelism](#): you can use `torch.nn.DataParallel` to wrap any module and it will be (almost magically) parallelized over batch dimension. This way you can leverage multiple GPUs with almost no effort.

On the other hand, TensorFlow allows you to fine tune every operation to be run on specific device. Nonetheless, defining parallelism is way more manual and requires careful thought. Consider the code that implements something like `DataParallel` in TensorFlow:

```
def make_parallel(fn, num_gpus, **kwargs):
    in_splits = {}
    for k, v in kwargs.items():
        in_splits[k] = tf.split(v, num_gpus)

    out_split = []
    for i in range(num_gpus):
        with tf.device(tf.DeviceSpec(device_type="GPU",
```

Get started

Open in app



```

        out_split.append((in_split * v[1] for k, v in
in_splits.items()))

return tf.concat(out_split, axis=0)

def model(a, b):
    return a + b

c = make_parallel(model, 2, a=a, b=b)

```

That being said, when using TensorFlow you can achieve everything you can do in PyTorch, but with more effort (you have more control as a bonus).

Also it is worth noting that both frameworks support distributed execution and provide high level interfaces for defining clusters.

Difference #6 — A Framework or a library

Let's build a CNN classifier for handwritten digits. Now PyTorch will really start to look like a *framework*. Recall that a programming framework gives us useful abstractions in certain domain and a convenient way to use them to solve concrete problems. That is the essence that separates a framework from a library.

Here we introduce `datasets` module which contains wrappers for popular datasets used to benchmark deep learning architectures. Also `nn.Module` is used to build a custom convolutional neural network classifier. `nn.Module` is a building block PyTorch gives us to create complex deep learning architectures. There are large amounts of ready to use modules in `torch.nn` package that we can use as a base for our model. Notice how PyTorch uses object oriented approach to define basic building blocks and give us some 'rails' to move on while providing ability to extend functionality via subclassing.

Here goes a slightly modified version of

<https://github.com/pytorch/examples/blob/master/mnist/main.py>:

```

1  import numpy as np
2  import tensorflow as tf
3
4  from tensorflow.contrib import learn

```

Get started

Open in app



```
8
9
10 def cnn_model_fn(features, labels, mode):
11     """Model function for CNN."""
12     # Input Layer
13     # Reshape X to 4-D tensor: [batch_size, width, height, channels]
14     # MNIST images are 28x28 pixels, and have one color channel
15     input_layer = tf.reshape(features, [-1, 28, 28, 1])
16
17     # Convolutional Layer #1
18     # Computes 32 features using a 5x5 filter with ReLU activation.
19     # Padding is added to preserve width and height.
20     # Input Tensor Shape: [batch_size, 28, 28, 1]
21     # Output Tensor Shape: [batch_size, 28, 28, 32]
22     conv1 = tf.layers.conv2d(
23         inputs=input_layer,
24         filters=32,
25         kernel_size=[5, 5],
26         padding="same",
27         activation=tf.nn.relu)
28
29     # Pooling Layer #1
30     # First max pooling layer with a 2x2 filter and stride of 2
31     # Input Tensor Shape: [batch_size, 28, 28, 32]
32     # Output Tensor Shape: [batch_size, 14, 14, 32]
33     pool1 = tf.layers.max_pooling2d(inputs=conv1, pool_size=[2, 2], strides=2)
34
35     # Convolutional Layer #2
36     # Computes 64 features using a 5x5 filter.
37     # Padding is added to preserve width and height.
38     # Input Tensor Shape: [batch_size, 14, 14, 32]
39     # Output Tensor Shape: [batch_size, 14, 14, 64]
40     conv2 = tf.layers.conv2d(
41         inputs=pool1,
42         filters=64,
43         kernel_size=[5, 5],
44         padding="same",
45         activation=tf.nn.relu)
46
47     # Pooling Layer #2
48     # Second max pooling layer with a 2x2 filter and stride of 2
49     # Input Tensor Shape: [batch_size, 14, 14, 64]
```

Get started

Open in app



```
52
53     # Flatten tensor into a batch of vectors
54     # Input Tensor Shape: [batch_size, 7, 7, 64]
55     # Output Tensor Shape: [batch_size, 7 * 7 * 64]
56     pool2_flat = tf.reshape(pool2, [-1, 7 * 7 * 64])
57
58     # Dense Layer
59     # Densely connected layer with 1024 neurons
60     # Input Tensor Shape: [batch_size, 7 * 7 * 64]
61     # Output Tensor Shape: [batch_size, 1024]
62     dense = tf.layers.dense(inputs=pool2_flat, units=1024, activation=tf.nn.relu)
63
64     # Add dropout operation; 0.6 probability that element will be kept
65     dropout = tf.layers.dropout(
66         inputs=dense, rate=0.4, training=mode == learn.ModeKeys.TRAIN)
67
68     # Logits layer
69     # Input Tensor Shape: [batch_size, 1024]
70     # Output Tensor Shape: [batch_size, 10]
71     logits = tf.layers.dense(inputs=dropout, units=10)
72
73     loss = None
74     train_op = None
75
76     # Calculate Loss (for both TRAIN and EVAL modes)
77     if mode != learn.ModeKeys.INFER:
78         onehot_labels = tf.one_hot(indices=tf.cast(labels, tf.int32), depth=10)
79         loss = tf.losses.softmax_cross_entropy(
80             onehot_labels=onehot_labels, logits=logits)
81
82     # Configure the Training Op (for TRAIN mode)
83     if mode == learn.ModeKeys.TRAIN:
84         train_op = tf.contrib.layers.optimize_loss(
85             loss=loss,
86             global_step=tf.contrib.framework.get_global_step(),
87             learning_rate=0.001,
88             optimizer="SGD")
89
90     # Generate Predictions
91     predictions = {
92         "classes": tf.argmax(
93             input=logits, axis=1),
```

[Get started](#)[Open in app](#)

```
97
98     # Return a ModelFnOps object
99     return model_fn_lib.ModelFnOps(
100         mode=mode, predictions=predictions, loss=loss, train_op=train_op)
101
102
103 # Load training and eval data
104 mnist = learn.datasets.load_dataset("mnist")
105 train_data = mnist.train.images # Returns np.array
106 train_labels = np.asarray(mnist.train.labels, dtype=np.int32)
107 eval_data = mnist.test.images # Returns np.array
108 eval_labels = np.asarray(mnist.test.labels, dtype=np.int32)
109
110 # Create the Estimator
111 mnist_classifier = learn.Estimator(
112     model_fn=cnn_model_fn, model_dir="/tmp/mnist_convnet_model")
113
114 # Set up logging for predictions
115 # Log the values in the "Softmax" tensor with label "probabilities"
116 tensors_to_log = {"probabilities": "softmax_tensor"}
117 logging_hook = tf.train.LoggingTensorHook(
118     tensors=tensors_to_log, every_n_iter=50)
119
120 # Train the model
121 mnist_classifier.fit(
122     x=train_data,
123     y=train_labels,
124     batch_size=100,
125     steps=20000,
126     monitors=[logging_hook])
127
128 # Configure the accuracy metric for evaluation
129 metrics = {
130     "accuracy":
131         learn.MetricSpec(
132             metric_fn=tf.metrics.accuracy, prediction_key="classes"),
133 }
134
135 # Evaluate the model and print results
136 eval_results = mnist_classifier.evaluate(
137     x=eval_data, y=eval_labels, metrics=metrics)
138 print(eval_results)
```


[Get started](#)[Open in app](#)

Plain TensorFlow feels a lot more like a library rather than a framework: all operations are pretty low-level and you will need to write lots of boilerplate code even when you might not want to (let's define those biases and weights again and again and ...).

As the time as passed a whole ecosystem of high-level wrappers started to emerge around TensorFlow. Each of those aims to simplify the way you work with the library. Many of them are currently located at `tensorflow.contrib` module (which is not considered a stable API) and some started to migrate to the main repository (see `tf.layers`).

So, you have a lot of freedom on how to use TensorFlow and what framework will suit the task best: [TFLearn](#), [tf.contrib.learn](#), [Sonnet](#), [Keras](#), plain `tf.layers`, etc. To be honest, Keras deserves another post but is currently out of the scope of this comparison.

Here we will use `tf.layers` and `tf.contrib.learn` to build our CNN classifier. The code follows the [official tutorial on tf.layers](#):

```
1  import numpy as np
2  import tensorflow as tf
3
4  from tensorflow.contrib import learn
5  from tensorflow.contrib.learn.python.learn.estimators import model_fn as model_fn_lib
6
7  tf.logging.set_verbosity(tf.logging.INFO)
8
9
10 def cnn_model_fn(features, labels, mode):
11     """Model function for CNN."""
12     # Input Layer
13     # Reshape X to 4-D tensor: [batch_size, width, height, channels]
14     # MNIST images are 28x28 pixels, and have one color channel
15     input_layer = tf.reshape(features, [-1, 28, 28, 1])
16
17     # Convolutional Layer #1
18     # Computes 32 features using a 5x5 filter with ReLU activation.
19     # Padding is added to preserve width and height.
20     # Input Tensor Shape: [batch_size, 28, 28, 1]
```

Get started

Open in app



```
24     filters=32,
25     kernel_size=[5, 5],
26     padding="same",
27     activation=tf.nn.relu)
28
29 # Pooling Layer #1
30 # First max pooling layer with a 2x2 filter and stride of 2
31 # Input Tensor Shape: [batch_size, 28, 28, 32]
32 # Output Tensor Shape: [batch_size, 14, 14, 32]
33 pool1 = tf.layers.max_pooling2d(inputs=conv1, pool_size=[2, 2], strides=2)
34
35 # Convolutional Layer #2
36 # Computes 64 features using a 5x5 filter.
37 # Padding is added to preserve width and height.
38 # Input Tensor Shape: [batch_size, 14, 14, 32]
39 # Output Tensor Shape: [batch_size, 14, 14, 64]
40 conv2 = tf.layers.conv2d(
41     inputs=pool1,
42     filters=64,
43     kernel_size=[5, 5],
44     padding="same",
45     activation=tf.nn.relu)
46
47 # Pooling Layer #2
48 # Second max pooling layer with a 2x2 filter and stride of 2
49 # Input Tensor Shape: [batch_size, 14, 14, 64]
50 # Output Tensor Shape: [batch_size, 7, 7, 64]
51 pool2 = tf.layers.max_pooling2d(inputs=conv2, pool_size=[2, 2], strides=2)
52
53 # Flatten tensor into a batch of vectors
54 # Input Tensor Shape: [batch_size, 7, 7, 64]
55 # Output Tensor Shape: [batch_size, 7 * 7 * 64]
56 pool2_flat = tf.reshape(pool2, [-1, 7 * 7 * 64])
57
58 # Dense Layer
59 # Densely connected layer with 1024 neurons
60 # Input Tensor Shape: [batch_size, 7 * 7 * 64]
61 # Output Tensor Shape: [batch_size, 1024]
62 dense = tf.layers.dense(inputs=pool2_flat, units=1024, activation=tf.nn.relu)
63
64 # Add dropout operation; 0.6 probability that element will be kept
65 dropout = tf.layers.dropout(
```

Get started

Open in app



```
68 # Logits layer
69 # Input Tensor Shape: [batch_size, 1024]
70 # Output Tensor Shape: [batch_size, 10]
71 logits = tf.layers.dense(inputs=dropout, units=10)
72
73 loss = None
74 train_op = None
75
76 # Calculate Loss (for both TRAIN and EVAL modes)
77 if mode != learn.ModeKeys.INFER:
78     onehot_labels = tf.one_hot(indices=tf.cast(labels, tf.int32), depth=10)
79     loss = tf.losses.softmax_cross_entropy(
80         onehot_labels=onehot_labels, logits=logits)
81
82 # Configure the Training Op (for TRAIN mode)
83 if mode == learn.ModeKeys.TRAIN:
84     train_op = tf.contrib.layers.optimize_loss(
85         loss=loss,
86         global_step=tf.contrib.framework.get_global_step(),
87         learning_rate=0.001,
88         optimizer="SGD")
89
90 # Generate Predictions
91 predictions = {
92     "classes": tf.argmax(
93         input=logits, axis=1),
94     "probabilities": tf.nn.softmax(
95         logits, name="softmax_tensor")
96 }
97
98 # Return a ModelFnOps object
99 return model_fn_lib.ModelFnOps(
100     mode=mode, predictions=predictions, loss=loss, train_op=train_op)
101
102
103 # Load training and eval data
104 mnist = learn.datasets.load_dataset("mnist")
105 train_data = mnist.train.images # Returns np.array
106 train_labels = np.asarray(mnist.train.labels, dtype=np.int32)
107 eval_data = mnist.test.images # Returns np.array
108 eval_labels = np.asarray(mnist.test.labels, dtype=np.int32)
109
```

[Get started](#)[Open in app](#)

```
113
114 # Set up logging for predictions
115 # Log the values in the "Softmax" tensor with label "probabilities"
116 tensors_to_log = {"probabilities": "softmax_tensor"}
117 logging_hook = tf.train.LoggingTensorHook(
118     tensors=tensors_to_log, every_n_iter=50)
119
120 # Train the model
121 mnist_classifier.fit(
122     x=train_data,
123     y=train_labels,
124     batch_size=100,
125     steps=20000,
126     monitors=[logging_hook])
127
128 # Configure the accuracy metric for evaluation
129 metrics = {
130     "accuracy":
131         learn.MetricSpec(
132             metric_fn=tf.metrics.accuracy, prediction_key="classes"),
133 }
134
135 # Evaluate the model and print results
136 eval_results = mnist_classifier.evaluate(
137     x=eval_data, y=eval_labels, metrics=metrics)
138 print(eval_results)
```

tensorflow_mnist.py hosted with ❤ by GitHub

[view raw](#)

So, both TensorFlow and PyTorch provide useful abstractions to reduce amounts of boilerplate code and speed up model development. The main difference between them is that PyTorch may feel more “pythonic” and has an object-oriented approach while TensorFlow has several options from which you may choose.

Personally, I consider PyTorch to be more clear and developer-friendly. It’s `torch.nn.Module` gives you the ability to define reusable modules in an OOP manner and I find this approach very flexible and powerful. Later you can compose all kind of modules via `torch.nn.Sequential` (hi Keras 🙌). Also, you have all built-in modules in a

[Get started](#)[Open in app](#)

Of course, you can write very clean code in plain TensorFlow but it just takes more skill and trial-and-error before you get it. When it goes to higher-level frameworks such as Keras or TFLearn get ready to lose at least some of the flexibility TensorFlow has to offer.

Conclusion

TensorFlow is very powerful and mature deep learning library with strong visualization capabilities and several options to use for high-level model development. It has production-ready deployment options and support for mobile platforms. TensorFlow is a good option if you:

- Develop models for production
- Develop models which need to be deployed on mobile platforms
- Want good community support and comprehensive documentation
- Want rich learning resources in various forms (TensorFlow has an entire [MOOC](#))
- Want or need to use Tensorboard
- Need to use large-scale distributed model training

PyTorch is still a young framework which is getting momentum fast. You may find it a good fit if you:

- Do research or your production non-functional requirements are not very demanding
- Want better development and debugging experience
- *Love* all things Pythonic

If you have the time the best advice would be to try both and see what fits your needs best.

If you liked this article, please leave a few 🙌. It lets me know that I am helping.

[Get started](#)[Open in app](#)

native C++ API, JIT compilation and ONNX integration. This means that you will be able to write production-ready services and do what TensorFlow Serving does. This is a big step to PyTorch and surely will empower its position as a fully featured framework for both research and production purposes.

Contact us

Need help with TensorFlow or PyTorch? Contact us at datalab@cinimex.ru

Want to get regular stream of interesting resources on Data Science and Entrepreneurship?

Follow me on  [Twitter](#) and  [Medium](#).

Sign up for The Variable

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)

Your email

[Get this newsletter](#)

By signing up, you will create a Medium account if you don't already have one. Review our [Privacy Policy](#) for more information about our privacy practices.

[Machine Learning](#)[Deep Learning](#)[TensorFlow](#)[Python](#)[Pytorch](#)[About](#) [Help](#) [Legal](#)